

output: html\_document: default pdf\_document: default — Title: "Classification" Authors: Gaurang Goel (GXG190015), Anurag Diwate (AXD190004) Date: 02/18/2023

**#Classification in general terms** Classification is a fundamental problem in machine learning in which the class of a new observation is predicted based on its attributes. Linear models for classification are a type of method that uses a linear function to represent the relationship between features and classes. In other words, these models learn a feature space linear boundary that divides the distinct classes.

**#Strengths and Weaknesses of Classification** One of the strengths of linear models for categorization is their ease of interpretation. They are simple to understand and depict, making them an excellent tool for exploratory data analysis. Linear models are also capable of dealing with high-dimensional data and scaling effectively to big datasets. They can also provide probabilistic forecasts, which may help in decision making.

Linear models, on the other hand, have some restrictions. They may not perform well when the feature-class connection is not linear, and they may fail to capture complicated relationships between features. Moreover, linear models are sensitive to outliers and may struggle with unbalanced classes. Nevertheless, these flaws are frequently overcome by employing more complex approaches such as feature engineering, regularization, or ensemble methods.

**#Data Source** <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud> (<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>)

```
#importing libraries
library(tidyverse)
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.4.1      ✓ purrr  1.0.1
## ✓ tibble  3.1.8      ✓ dplyr  1.1.0
## ✓ tidyr   1.3.0      ✓ stringr 1.5.0
## ✓ readr   2.1.4      ✓ forcats 1.0.0
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
## lift
```

```
# Importing dataset
data <- read.csv("e:/creditcard.csv", header = TRUE)

#attaching data
attach(data)

#Part A: Dividing the data into 80/20 train/test
set.seed(123)
trainIndex <- createDataPartition(data$Class, p = 0.8, list = FALSE)
train <- data[trainIndex, ]
test <- data[-trainIndex, ]

#Part B: Data Exploration
#names() method returns the names of the headers in the dataset
names(train)
```

```
## [1] "Time"    "V1"      "V2"      "V3"      "V4"      "V5"      "V6"      "V7"
## [9] "V8"      "V9"      "V10"     "V11"     "V12"     "V13"     "V14"     "V15"
## [17] "V16"     "V17"     "V18"     "V19"     "V20"     "V21"     "V22"     "V23"
## [25] "V24"     "V25"     "V26"     "V27"     "V28"     "Amount"  "Class"
```

```
#nrow() method returns the number of rows in the dataset
nrow(train)
```

```
## [1] 227846
```

```
#nrow() method returns the number of columns in the dataset
ncol(train)
```

```
## [1] 31
```

```
#The summary() method provides an overview of each variable's distribution in the dataset.
summary(train)
```

##	Time	V1	V2	V3
##	Min. : 0	Min. : -56.40751	Min. : -72.71573	Min. : -48.32559
##	1st Qu.: 54245	1st Qu.: -0.92091	1st Qu.: -0.59684	1st Qu.: -0.89049
##	Median : 84788	Median : 0.01692	Median : 0.06659	Median : 0.17970
##	Mean : 94865	Mean : -0.00086	Mean : 0.00174	Mean : 0.00027
##	3rd Qu.: 139370	3rd Qu.: 1.31581	3rd Qu.: 0.80442	3rd Qu.: 1.02775
##	Max. : 172792	Max. : 2.45189	Max. : 22.05773	Max. : 9.38256
##	V4	V5	V6	
##	Min. : -5.683171	Min. : -113.74331	Min. : -26.16051	
##	1st Qu.: -0.850256	1st Qu.: -0.69242	1st Qu.: -0.76871	
##	Median : -0.022409	Median : -0.05594	Median : -0.27562	
##	Mean : -0.002627	Mean : -0.00115	Mean : -0.00063	
##	3rd Qu.: 0.739540	3rd Qu.: 0.60994	3rd Qu.: 0.39724	
##	Max. : 16.875344	Max. : 34.80167	Max. : 73.30163	
##	V7	V8	V9	
##	Min. : -43.55724	Min. : -73.21672	Min. : -13.434066	
##	1st Qu.: -0.55464	1st Qu.: -0.20861	1st Qu.: -0.642508	
##	Median : 0.04025	Median : 0.02212	Median : -0.051192	
##	Mean : -0.00100	Mean : -0.00195	Mean : 0.000182	
##	3rd Qu.: 0.57005	3rd Qu.: 0.32705	3rd Qu.: 0.597453	
##	Max. : 120.58949	Max. : 20.00721	Max. : 15.594995	
##	V10	V11	V12	
##	Min. : -24.588262	Min. : -4.797473	Min. : -18.683715	
##	1st Qu.: -0.536870	1st Qu.: -0.765546	1st Qu.: -0.404669	
##	Median : -0.092855	Median : -0.035046	Median : 0.140933	
##	Mean : -0.000605	Mean : -0.001782	Mean : -0.000091	
##	3rd Qu.: 0.454017	3rd Qu.: 0.739648	3rd Qu.: 0.617885	
##	Max. : 23.745136	Max. : 12.018913	Max. : 7.848392	
##	V13	V14	V15	
##	Min. : -5.791881	Min. : -19.214325	Min. : -4.498945	
##	1st Qu.: -0.647447	1st Qu.: -0.426030	1st Qu.: -0.583692	
##	Median : -0.012773	Median : 0.051201	Median : 0.047257	
##	Mean : 0.001049	Mean : -0.000157	Mean : -0.000627	
##	3rd Qu.: 0.663943	3rd Qu.: 0.492691	3rd Qu.: 0.649871	
##	Max. : 7.126883	Max. : 10.526766	Max. : 8.877742	
##	V16	V17	V18	
##	Min. : -13.563273	Min. : -25.16280	Min. : -9.498746	
##	1st Qu.: -0.468480	1st Qu.: -0.48332	1st Qu.: -0.499053	
##	Median : 0.065990	Median : -0.06591	Median : -0.004552	
##	Mean : -0.000146	Mean : -0.00046	Mean : -0.000341	
##	3rd Qu.: 0.522553	3rd Qu.: 0.39914	3rd Qu.: 0.499546	
##	Max. : 17.315112	Max. : 9.25353	Max. : 5.041069	
##	V19	V20	V21	
##	Min. : -7.213527	Min. : -54.49772	Min. : -34.83038	
##	1st Qu.: -0.456252	1st Qu.: -0.21134	1st Qu.: -0.22855	
##	Median : 0.002498	Median : -0.06267	Median : -0.02982	
##	Mean : -0.000934	Mean : -0.00056	Mean : -0.00103	
##	3rd Qu.: 0.456781	3rd Qu.: 0.13244	3rd Qu.: 0.18611	
##	Max. : 5.591971	Max. : 39.42090	Max. : 27.20284	
##	V22	V23	V24	
##	Min. : -10.933144	Min. : -44.80774	Min. : -2.83663	
##	1st Qu.: -0.542962	1st Qu.: -0.16202	1st Qu.: -0.35459	

```
## Median : 0.005536 Median : -0.01118 Median : 0.04074
## Mean : -0.000237 Mean : 0.00018 Mean : -0.00009
## 3rd Qu.: 0.529175 3rd Qu.: 0.14761 3rd Qu.: 0.44074
## Max. : 10.503090 Max. : 22.08354 Max. : 4.58455
## V25 V26 V27
## Min. : -10.295397 Min. : -2.604551 Min. : -22.565679
## 1st Qu.: -0.316603 1st Qu.: -0.326462 1st Qu.: -0.070855
## Median : 0.017095 Median : -0.051426 Median : 0.001261
## Mean : 0.000646 Mean : 0.000532 Mean : -0.000043
## 3rd Qu.: 0.351245 3rd Qu.: 0.241546 3rd Qu.: 0.091068
## Max. : 7.519589 Max. : 3.517346 Max. : 31.612198
## V28 Amount Class
## Min. : -15.43008 Min. : 0.00 Min. : 0.00000
## 1st Qu.: -0.05293 1st Qu.: 5.60 1st Qu.: 0.00000
## Median : 0.01128 Median : 22.00 Median : 0.00000
## Mean : 0.00053 Mean : 88.16 Mean : 0.00169
## 3rd Qu.: 0.07846 3rd Qu.: 77.05 3rd Qu.: 0.00000
## Max. : 33.84781 Max. : 25691.16 Max. : 1.00000
```


*#The str() method displays the dataset's structure, including variable data types.*  
 str(train)

```
## 'data.frame':  227846 obs. of  31 variables:
## $ Time   : num  0 0 1 1 4 7 7 9 10 10 ...
## $ V1     : num  -1.36 1.192 -1.358 -0.966 1.23 ...
## $ V2     : num  -0.0728 0.2662 -1.3402 -0.1852 0.141 ...
## $ V3     : num  2.5363 0.1665 1.7732 1.793 0.0454 ...
## $ V4     : num  1.378 0.448 0.38 -0.863 1.203 ...
## $ V5     : num  -0.3383 0.06 -0.5032 -0.0103 0.1919 ...
## $ V6     : num  0.4624 -0.0824 1.8005 1.2472 0.2727 ...
## $ V7     : num  0.2396 -0.0788 0.79146 0.23761 -0.00516 ...
## $ V8     : num  0.0987 0.0851 0.2477 0.3774 0.0812 ...
## $ V9     : num  0.364 -0.255 -1.515 -1.387 0.465 ...
## $ V10    : num  0.0908 -0.167 0.2076 -0.055 -0.0993 ...
## $ V11    : num  -0.552 1.613 0.625 -0.226 -1.417 ...
## $ V12    : num  -0.6178 1.0652 0.0661 0.1782 -0.1538 ...
## $ V13    : num  -0.991 0.489 0.717 0.508 -0.751 ...
## $ V14    : num  -0.311 -0.144 -0.166 -0.288 0.167 ...
## $ V15    : num  1.4682 0.6356 2.3459 -0.6314 0.0501 ...
## $ V16    : num  -0.47 0.464 -2.89 -1.06 -0.444 ...
## $ V17    : num  0.20797 -0.1148 1.10997 -0.68409 0.00282 ...
## $ V18    : num  0.0258 -0.1834 -0.1214 1.9658 -0.612 ...
## $ V19    : num  0.404 -0.1458 -2.2619 -1.2326 -0.0456 ...
## $ V20    : num  0.2514 -0.0691 0.525 -0.208 -0.2196 ...
## $ V21    : num  -0.0183 -0.2258 0.248 -0.1083 -0.1677 ...
## $ V22    : num  0.27784 -0.63867 0.77168 0.00527 -0.27071 ...
## $ V23    : num  -0.11 0.101 0.909 -0.19 -0.154 ...
## $ V24    : num  0.0669 -0.3398 -0.6893 -1.1756 -0.7801 ...
## $ V25    : num  0.129 0.167 -0.328 0.647 0.75 ...
## $ V26    : num  -0.189 0.126 -0.139 -0.222 -0.257 ...
## $ V27    : num  0.13356 -0.00898 -0.05535 0.06272 0.03451 ...
## $ V28    : num  -0.02105 0.01472 -0.05975 0.06146 0.00517 ...
## $ Amount: num  149.62 2.69 378.66 123.5 4.99 ...
## $ Class  : int  0 0 0 0 0 0 0 0 0 0 ...
```


### *#Part C: Graphs*

#### *# Density plot*

```
ggplot(train, aes(Amount, fill = factor(Class))) +
  geom_density(alpha = 0.5) +
  labs(x = "Amount", y = "Density", fill = "Class")
```



```
# Scatter plot of V1 and V2
ggplot(train, aes(V1, V2, color = factor(Class))) +
  geom_point(alpha = 0.5) +
  labs(x = "V1", y = "V2", color = "Class")
```



```
#Part D: Logistical Regression Model
```

```
log_model <- glm(Class ~ ., data = train, family = "binomial")
```

```
summary(log_model)
```

```
##
## Call:
## glm(formula = Class ~ ., family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.6403  -0.0285  -0.0189  -0.0122   4.3118
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -8.470e+00  2.807e-01 -30.177 < 2e-16 ***
## Time        -3.211e-06  2.552e-06  -1.258  0.20839
## V1           1.233e-01  4.839e-02   2.547  0.01086 *
## V2          -1.550e-02  6.446e-02  -0.241  0.80994
## V3           2.654e-03  6.100e-02   0.044  0.96530
## V4           6.894e-01  7.712e-02   8.939 < 2e-16 ***
## V5           5.457e-02  7.465e-02   0.731  0.46475
## V6          -1.392e-01  8.717e-02  -1.597  0.11029
## V7          -1.457e-01  7.483e-02  -1.947  0.05153 .
## V8          -1.725e-01  3.521e-02  -4.901 9.55e-07 ***
## V9          -3.556e-01  1.168e-01  -3.045  0.00233 **
## V10         -7.345e-01  1.019e-01  -7.211 5.56e-13 ***
## V11         -1.823e-02  9.165e-02  -0.199  0.84231
## V12          1.320e-01  9.646e-02   1.368  0.17117
## V13         -2.744e-01  9.289e-02  -2.954  0.00314 **
## V14         -5.992e-01  6.869e-02  -8.724 < 2e-16 ***
## V15         -1.288e-01  9.719e-02  -1.325  0.18508
## V16         -2.016e-01  1.305e-01  -1.545  0.12242
## V17          6.154e-02  7.624e-02   0.807  0.41957
## V18         -4.317e-02  1.358e-01  -0.318  0.75062
## V19          8.628e-02  1.058e-01   0.816  0.41478
## V20         -4.668e-01  9.268e-02  -5.037 4.74e-07 ***
## V21          3.860e-01  6.701e-02   5.760 8.40e-09 ***
## V22          6.102e-01  1.506e-01   4.051 5.11e-05 ***
## V23         -1.279e-01  6.830e-02  -1.872  0.06118 .
## V24          1.507e-01  1.675e-01   0.900  0.36821
## V25         -2.805e-02  1.491e-01  -0.188  0.85073
## V26          6.545e-02  2.126e-01   0.308  0.75818
## V27         -8.302e-01  1.360e-01  -6.105 1.03e-09 ***
## V28         -2.883e-01  1.022e-01  -2.822  0.00477 **
## Amount       9.044e-04  4.506e-04   2.007  0.04472 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 5684.4  on 227845  degrees of freedom
## Residual deviance: 1754.3  on 227815  degrees of freedom
## AIC: 1816.3
##
## Number of Fisher Scoring iterations: 11
```



*#The summary contains metrics for evaluating the model, such as deviation, AIC, and likelihood ratio test. The AIC assesses the trade-off between model complexity and goodness of fit, whereas the deviation reflects the gap between anticipated and actual values. The likelihood ratio test compares the current model to a null model with no predictors and returns a p-value showing whether the current model is significantly better than the null model. Overall, the model summary is an excellent tool for assessing the logistic regression model and suggesting areas for improvement.*

*#Part E: Naïve Bayes Model*

```
library(e1071)  
nb_model <- naiveBayes(Class ~ ., data = train)  
nb_model
```

```
##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##           0           1
## 0.998310262 0.001689738
##
## Conditional probabilities:
##   Time
## Y      [,1]      [,2]
## 0 94890.48 47507.65
## 1 79860.10 48184.32
##
##   V1
## Y      [,1]      [,2]
## 0  0.00740885 1.932901
## 1 -4.88639316 7.128115
##
##   V2
## Y      [,1]      [,2]
## 0 -0.004798741 1.641011
## 1  3.867319951 4.367791
##
##   V3
## Y      [,1]      [,2]
## 0  0.01256929 1.460488
## 1 -7.26811935 7.325462
##
##   V4
## Y      [,1]      [,2]
## 0 -0.01046061 1.399779
## 1  4.62538083 2.868327
##
##   V5
## Y      [,1]      [,2]
## 0  0.004503653 1.363200
## 1 -3.338808334 5.580357
##
##   V6
## Y      [,1]      [,2]
## 0  0.001747588 1.332378
## 1 -1.407842835 1.870314
##
##   V7
## Y      [,1]      [,2]
## 0  0.008840788 1.187781
## 1 -5.813152477 7.584867
##
```

```
##      V8
## Y      [,1]      [,2]
## 0 -0.002771184 1.176258
## 1  0.481534185 7.190409
##
##      V9
## Y      [,1]      [,2]
## 0  0.004688844 1.091160
## 1 -2.662470701 2.596888
##
##      V10
## Y      [,1]      [,2]
## 0  0.009246005 1.044319
## 1 -5.820703369 5.102794
##
##      V11
## Y      [,1]      [,2]
## 0 -0.008330652 1.003939
## 1  3.867088122 2.643666
##
##      V12
## Y      [,1]      [,2]
## 0  0.0106774 0.9477549
## 1 -6.3618687 4.6697379
##
##      V13
## Y      [,1]      [,2]
## 0  0.001223421 0.9940628
## 1 -0.102097938 1.0828370
##
##      V14
## Y      [,1]      [,2]
## 0  0.01191161 0.8977033
## 1 -7.13024005 4.3136176
##
##      V15
## Y      [,1]      [,2]
## 0 -0.0004786208 0.9154533
## 1 -0.0884254463 1.0552750
##
##      V16
## Y      [,1]      [,2]
## 0  0.006911794 0.8452431
## 1 -4.170237462 3.8995826
##
##      V17
## Y      [,1]      [,2]
## 0  0.01095445 0.7509774
## 1 -6.74433901 7.1226709
##
##      V18
## Y      [,1]      [,2]
```

```
## 0 0.003518636 0.8251347
## 1 -2.280731156 2.9367788
##
## V19
## Y      [,1]      [,2]
## 0 -0.002026151 0.8107966
## 1 0.644302710 1.5437516
##
## V20
## Y      [,1]      [,2]
## 0 -0.001244412 0.7792088
## 1 0.402485116 1.3860709
##
## V21
## Y      [,1]      [,2]
## 0 -0.002178711 0.7210648
## 1 0.680381563 4.1045657
##
## V22
## Y      [,1]      [,2]
## 0 -0.0002308207 0.7238815
## 1 -0.0038244612 1.5596719
##
## V23
## Y      [,1]      [,2]
## 0 0.0002707705 0.6185419
## 1 -0.0560149139 1.7222627
##
## V24
## Y      [,1]      [,2]
## 0 7.543347e-05 0.6061957
## 1 -9.774902e-02 0.5212982
##
## V25
## Y      [,1]      [,2]
## 0 0.0005272097 0.5214561
## 1 0.0708755337 0.8202603
##
## V26
## Y      [,1]      [,2]
## 0 0.0004257687 0.4825454
## 1 0.0631377084 0.4887205
##
## V27
## Y      [,1]      [,2]
## 0 -0.0002412351 0.402310
## 1 0.1169885093 1.486415
##
## V28
## Y      [,1]      [,2]
## 0 0.0003755148 0.3368746
## 1 0.0916224446 0.5324919
```

```
##
##   Amount
## Y      [,1]      [,2]
## 0  88.11217 252.6090
## 1 117.78574 244.7329
```

*#Based on the training data, the Naive Bayes model learns the conditional probability of each characteristic for each class (fraudulent or non-fraudulent transaction). Specifically, the model evaluated the likelihood of a particular feature value appearing in a fraudulent or non-fraudulent transaction and utilized these probabilities to categorize fresh transactions in the test data. Naive Bayes is a probabilistic model that implies independence between characteristics given a class, resulting in a simpler model and faster training times than other models. This assumption, however, may not always hold true in practice, and the model may suffer from underfitting if crucial relationships are not recorded.*

*#Part F: Predict and Evaluate*

*# Logistic regression predictions*

```
log_pred <- predict(log_model, newdata = test, type = "response")
log_pred_class <- ifelse(log_pred > 0.5, 1, 0)
```

*# Naïve Bayes predictions*

```
nb_pred <- predict(nb_model, newdata = test)
nb_pred_class <- as.numeric(as.character(nb_pred)) - 1
```

*# Logistic Regression Metrics*

```
log_confusion <- confusionMatrix(as.factor(log_pred_class), as.factor(test$Class))
log_accuracy <- log_confusion$overall[1]
log_sensitivity <- log_confusion$byClass[1]
log_specificity <- log_confusion$byClass[2]
log_precision <- log_confusion$byClass[3]
log_f1_score <- log_confusion$byClass[4]
```

*# Naive Bayes Metrics*

```
nb_pred_class <- factor(nb_pred_class, levels = c("0", "1"))
nb_confusion <- confusionMatrix(as.factor(nb_pred_class), as.factor(test$Class))
nb_accuracy <- nb_confusion$overall[1]
nb_sensitivity <- nb_confusion$byClass[1]
nb_specificity <- nb_confusion$byClass[2]
nb_precision <- nb_confusion$byClass[3]
nb_f1_score <- nb_confusion$byClass[4]
if (is.nan(nb_f1_score)) {
  nb_f1_score <- 0
}
```

*# Print metrics*

```
print(paste0("Logistic Regression Accuracy: ", round(log_accuracy, 3)))
```

```
## [1] "Logistic Regression Accuracy: 0.999"
```

```
print(paste0("Logistic Regression Sensitivity: ", round(log_sensitivity, 3)))
```

```
## [1] "Logistic Regression Sensitivity: 1"
```

```
print(paste0("Logistic Regression Specificity: ", round(log_specificity, 3)))
```

```
## [1] "Logistic Regression Specificity: 0.579"
```

```
print(paste0("Logistic Regression Precision: ", round(log_precision, 3)))
```

```
## [1] "Logistic Regression Precision: 0.999"
```

```
print(paste0("Logistic Regression F1 Score: ", round(log_f1_score, 3)))
```

```
## [1] "Logistic Regression F1 Score: 0.861"
```

```
print(paste0("Naive Bayes Accuracy: ", round(nb_accuracy, 3)))
```

```
## [1] "Naive Bayes Accuracy: 0.932"
```

```
print(paste0("Naive Bayes Sensitivity: ", round(nb_sensitivity, 3)))
```

```
## [1] "Naive Bayes Sensitivity: 1"
```

```
print(paste0("Naive Bayes Specificity: ", round(nb_specificity, 3)))
```

```
## [1] "Naive Bayes Specificity: 0"
```

```
print(paste0("Naive Bayes Precision: ", round(nb_precision, 3)))
```

```
## [1] "Naive Bayes Precision: 0.932"
```

```
print(paste0("Naive Bayes F1 Score: ", round(nb_f1_score, 3)))
```

```
## [1] "Naive Bayes F1 Score: 0"
```

**#Part G: Strengths and Weaknesses of Naïve Bayes and Logistic Regression** In machine learning, two common algorithms for classification problems are Nave Bayes and Logistic Regression.

Nave Bayes is a probabilistic method that performs well with large datasets and requires less training data than other algorithms. It is simple and efficient, making it an excellent choice for real-time applications. Nevertheless, in real-world datasets, Nave Bayes presupposes that all characteristics are independent of one another. When dealing with associated characteristics, this might lead to erroneous forecasts.

Logistic Regression is a linear approach that is effective for binary classification issues. It is simple to use and understand, making it a popular choice for both beginners and professionals. Logistic Regression also generates probabilistic outputs that can be used to make decisions. Nevertheless, when dealing with complicated datasets, Logistic Regression presupposes a linear relationship between the independent factors and the log-odds of the dependent variable, which can lead to underfitting. Moreover, Logistic Regression might struggle with multicollinear datasets, which include linked independent variables.

#Part H: Benefits and Drawbacks of the classification metrics used  
Accuracy: Benefit: It is simple and easy to grasp, and it gauges the overall accuracy of the model predictions. Drawback: If the classes are uneven, great accuracy can be obtained simply by forecasting the majority class.

Sensitivity: Benefit: Measures the model's ability to properly identify positive situations, which is useful in applications where the cost of false negatives is large (e.g. medical diagnoses) Drawback: May be less essential in cases where the cost of false positives is considerable.

Specificity: Benefit: Indicates the model's ability to properly identify negative situations, which is useful in applications where the cost of false positives is large (e.g. fraud detection) Drawback: May be less essential in applications where the cost of false negatives is considerable.

Precision: Benefit: Determines the fraction of real positive predictions among all positive predictions, which is useful in situations where the cost of false positives is large (e.g. spam email detection) Drawback: May be less essential in applications where the cost of false negatives is considerable.

F1 Score: Benefit: Combines accuracy and recall into a single statistic, which can be advantageous for optimizing for both metrics simultaneously. Drawback: Because it is a weighted average of the two measurements, it may be less interpretable than the individual metrics.